

a) Apply Decision Tree, Random Forest, and KNN

```
# Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score,
f1_score, precision_score, recall_score, roc_curve, auc
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset
data = pd.read_csv(r"C:\Users\bolla\Downloads\msba sem1\scm\Prostate
Cancer.csv")
```

```
data.head()
```

	Gender	Smoking	Prostate	Volume	PSA	Level	Age	Prostate	Cancer
0	Male	No		50.89		9.64	57		No
1	Female	No		27.95		2.89	41		No
2	Male	Yes		20.22		5.22	74		No
3	Male	No		52.62		3.36	55		No
4	Male	Yes		48.27		3.21	80		No

```
data.describe()
```

	Prostate Volume	PSA Level	Age
count	70.000000	70.000000	70.000000
mean	39.675571	4.975714	66.557143
std	11.716884	2.844250	13.562134
min	20.220000	0.550000	41.000000
25%	29.355000	2.677500	57.250000
50%	40.045000	5.250000	69.000000
75%	50.737500	7.030000	74.750000
max	58.870000	9.860000	88.000000

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 70 entries, 0 to 69
Data columns (total 6 columns):
 #   Column              Non-Null Count  Dtype
 0   ...                  ...              ...
 5   ...                  ...              ...
 6   ...                  ...              ...
 7   ...                  ...              ...
 8   ...                  ...              ...
 9   ...                  ...              ...
10   ...                  ...              ...
11   ...                  ...              ...
12   ...                  ...              ...
13   ...                  ...              ...
14   ...                  ...              ...
15   ...                  ...              ...
16   ...                  ...              ...
17   ...                  ...              ...
18   ...                  ...              ...
19   ...                  ...              ...
20   ...                  ...              ...
21   ...                  ...              ...
22   ...                  ...              ...
23   ...                  ...              ...
24   ...                  ...              ...
25   ...                  ...              ...
26   ...                  ...              ...
27   ...                  ...              ...
28   ...                  ...              ...
29   ...                  ...              ...
30   ...                  ...              ...
31   ...                  ...              ...
32   ...                  ...              ...
33   ...                  ...              ...
34   ...                  ...              ...
35   ...                  ...              ...
36   ...                  ...              ...
37   ...                  ...              ...
38   ...                  ...              ...
39   ...                  ...              ...
40   ...                  ...              ...
41   ...                  ...              ...
42   ...                  ...              ...
43   ...                  ...              ...
44   ...                  ...              ...
45   ...                  ...              ...
46   ...                  ...              ...
47   ...                  ...              ...
48   ...                  ...              ...
49   ...                  ...              ...
50   ...                  ...              ...
51   ...                  ...              ...
52   ...                  ...              ...
53   ...                  ...              ...
54   ...                  ...              ...
55   ...                  ...              ...
56   ...                  ...              ...
57   ...                  ...              ...
58   ...                  ...              ...
59   ...                  ...              ...
60   ...                  ...              ...
61   ...                  ...              ...
62   ...                  ...              ...
63   ...                  ...              ...
64   ...                  ...              ...
65   ...                  ...              ...
66   ...                  ...              ...
67   ...                  ...              ...
68   ...                  ...              ...
69   ...                  ...              ...
```

```

---      -----      -----      -----
0  Gender          70 non-null    object
1  Smoking         70 non-null    object
2  Prostate Volume 70 non-null    float64
3  PSA Level       70 non-null    float64
4  Age            70 non-null    int64
5  Prostate Cancer 70 non-null    object
dtypes: float64(2), int64(1), object(3)
memory usage: 3.4+ KB

data.isnull().sum()# there is no missing values in the data set

Gender          0
Smoking         0
Prostate Volume 0
PSA Level       0
Age            0
Prostate Cancer 0
dtype: int64

def remove_outliers(df):
    columns_to_check = ['Prostate Volume', 'PSA Level', 'Age']
    for column in columns_to_check:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        df = df[(df[column] >= (Q1 - 1.5 * IQR)) & (df[column] <= (Q3
+ 1.5 * IQR))]
    return df

# Visualize the data before cleaning
plt.figure(figsize=(10, 5))
sns.boxplot(data=data[['Prostate Volume', 'PSA Level', 'Age']])
plt.title('Boxplot Before Data Cleaning (Outliers Present)')
plt.show()

data_cleaned = remove_outliers(data)

# Visualize the data after cleaning
plt.figure(figsize=(10, 5))
sns.boxplot(data=data_cleaned[['Prostate Volume', 'PSA Level',
'Age']])
plt.title('Boxplot After Data Cleaning (Outliers Removed)')
plt.show()

# Encode categorical variables
encoder = LabelEncoder()
data_cleaned['Gender'] = encoder.fit_transform(data_cleaned['Gender'])

data_cleaned['Smoking'] =

```

```

encoder.fit_transform(data_cleaned['Smoking'])
data_cleaned['Prostate Cancer'] =
encoder.fit_transform(data_cleaned['Prostate Cancer'])

# Split the dataset into features and target variable
X = data_cleaned.drop('Prostate Cancer', axis=1)
y = data_cleaned['Prostate Cancer']

# Visualize data distribution before balancing
plt.figure(figsize=(8, 5))
sns.countplot(x=y)
plt.title('Data Distribution Before Balancing')
plt.xlabel('Prostate Cancer (0: No, 1: Yes)')
plt.ylabel('Count')
plt.show()

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Apply SMOTE to the training data
smote = SMOTE(random_state=42)
X_train_balanced, y_train_balanced = smote.fit_resample(X_train,
y_train)

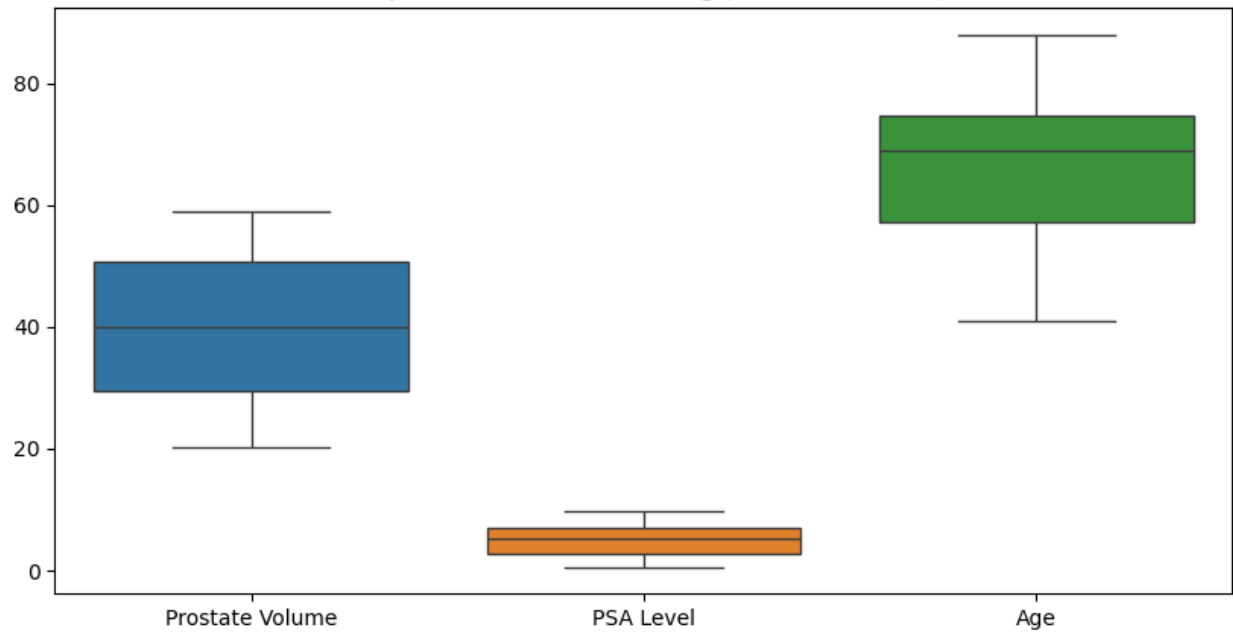
# Visualize the data after SMOTE
plt.figure(figsize=(8, 5))
sns.countplot(x=y_train_balanced)
plt.title('Data Distribution After SMOTE (Balanced)')
plt.xlabel('Prostate Cancer (0: No, 1: Yes)')
plt.ylabel('Count')
plt.show()

# Normalize the data for KNN
scaler = StandardScaler()
X_train_balanced = scaler.fit_transform(X_train_balanced)
X_test = scaler.transform(X_test)


# Visualize the data after scaling
plt.figure(figsize=(10, 5))
sns.boxplot(data=X_train_balanced)
plt.title('Boxplot After Data Normalization')
plt.xticks(ticks=np.arange(len(X.columns)), labels=X.columns,
rotation=45)
plt.show()

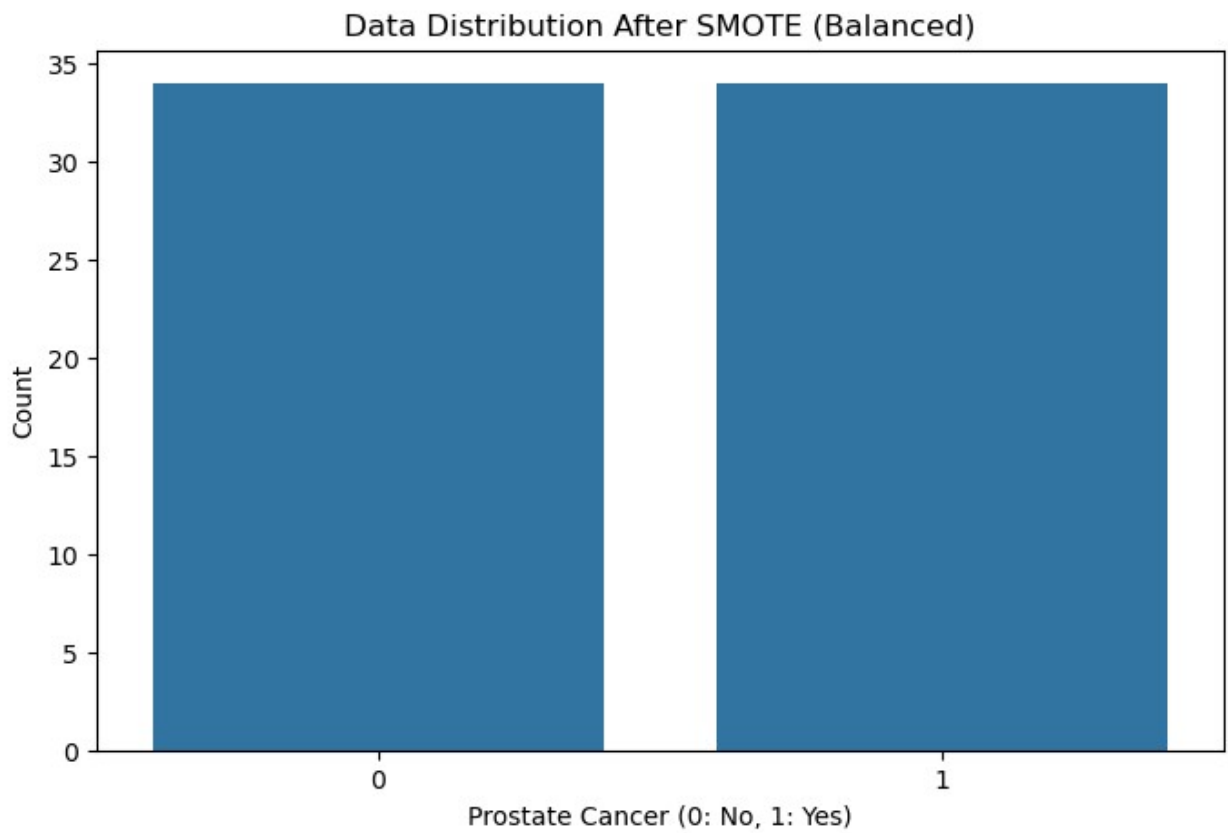
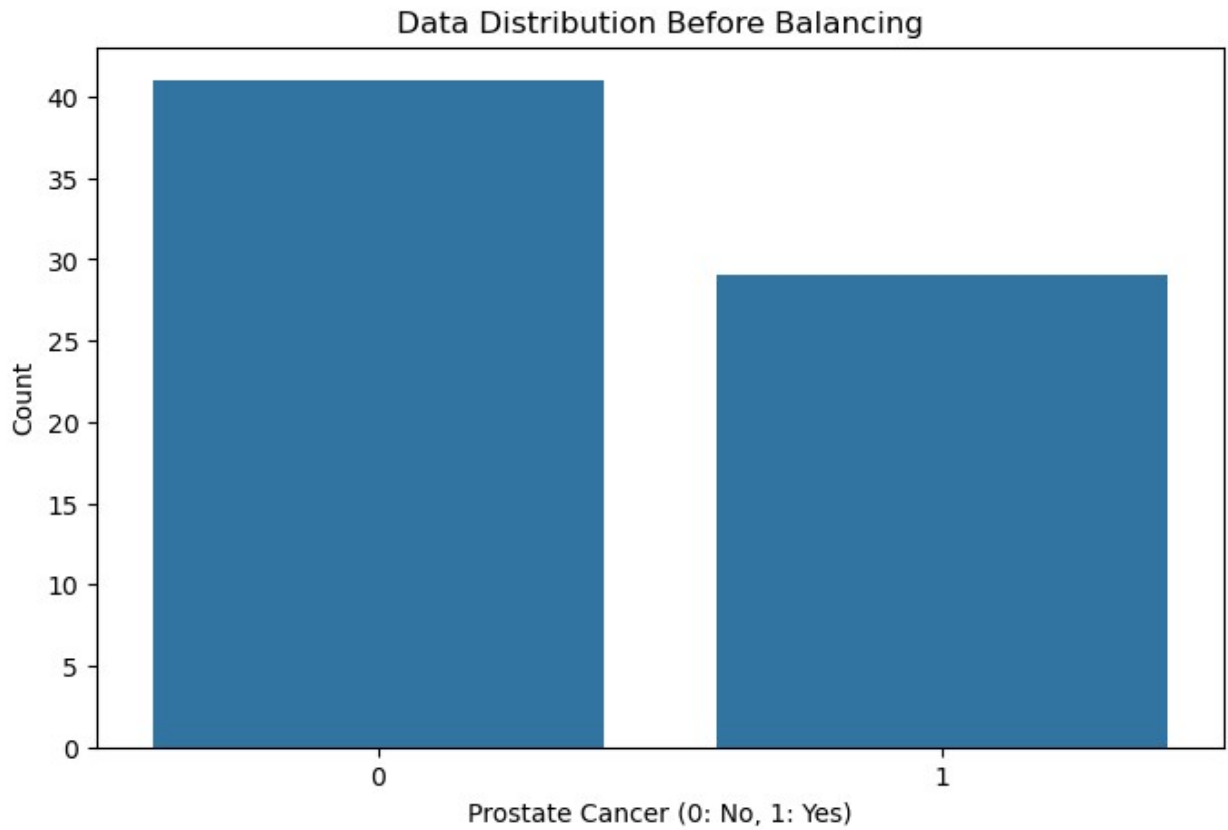
```

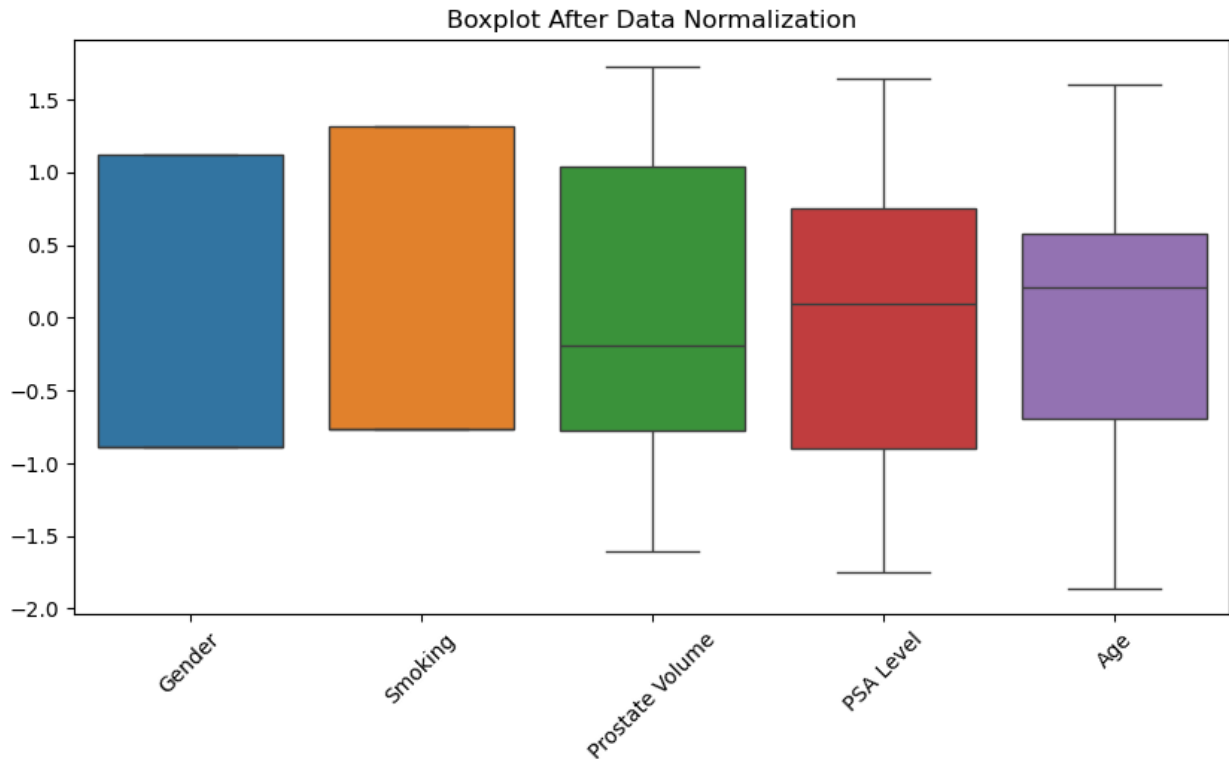
Boxplot Before Data Cleaning (Outliers Present)



Boxplot After Data Cleaning (Outliers Removed)







b) Show the confusion matrix, accuracy, F1 score, precision, and recall for training and testing datasets

```
# Initialize models
decision_tree = DecisionTreeClassifier(random_state=42)
random_forest = RandomForestClassifier(n_estimators=100,
criterion='gini', random_state=42)
knn = KNeighborsClassifier()

# Train models on balanced data
decision_tree.fit(X_train_balanced, y_train_balanced)
random_forest.fit(X_train_balanced, y_train_balanced)
knn.fit(X_train_balanced, y_train_balanced)

# Define a function to calculate and display model metrics
def display_metrics(model, X_train, y_train, X_test, y_test,
model_name):
    # Predictions
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)

    # Metrics
```

```

metrics = {
    'Accuracy': [accuracy_score(y_train, y_train_pred),
accuracy_score(y_test, y_test_pred)],
    'F1 Score': [f1_score(y_train, y_train_pred), f1_score(y_test,
y_test_pred)],
    'Precision': [precision_score(y_train, y_train_pred),
precision_score(y_test, y_test_pred)],
    'Recall': [recall_score(y_train, y_train_pred),
recall_score(y_test, y_test_pred)],
}

print(f'\n{model_name} Performance:')
for metric, values in metrics.items():
    print(f'{metric}: Train = {values[0]:.2f}, Test =
{values[1]:.2f}')

# Confusion Matrix
cm = confusion_matrix(y_test, y_test_pred)
plt.figure(figsize=(5, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Cancer', 'Cancer'],
            yticklabels=['No Cancer', 'Cancer'])
plt.title(f'Confusion Matrix for {model_name}')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Display metrics for all models
display_metrics(decision_tree, X_train_balanced, y_train_balanced,
X_test, y_test, 'Decision Tree')
display_metrics(random_forest, X_train_balanced, y_train_balanced,
X_test, y_test, 'Random Forest')
display_metrics(knn, X_train_balanced, y_train_balanced, X_test,
y_test, 'KNN')

# Finding the optimal k value for KNN
k_values = range(1, 21)
accuracy_scores = []

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_balanced, y_train_balanced)
    score = knn.score(X_test, y_test)
    accuracy_scores.append(score)

# Plotting accuracy for different k values
plt.figure(figsize=(10, 5))
plt.plot(k_values, accuracy_scores, marker='o')
plt.title('KNN Accuracy vs. K Values')
plt.xlabel('K Value')

```

```

plt.ylabel('Accuracy')
plt.xticks(k_values)
plt.grid()
plt.show()


# Identify the optimal k
optimal_k = k_values[np.argmax(accuracy_scores)]
print(f'The optimal k value for KNN is: {optimal_k}')

# Re-train the KNN model using the optimal k value
knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k)
knn_optimal.fit(X_train_balanced, y_train_balanced)


# Evaluate the model performance again using the testing dataset for KNN
display_metrics(knn_optimal, X_train_balanced, y_train_balanced,
X_test, y_test, 'KNN (Optimal)')

```



Decision Tree Performance:
Accuracy: Train = 1.00, Test = 0.71
F1 Score: Train = 1.00, Test = 0.67
Precision: Train = 1.00, Test = 0.80
Recall: Train = 1.00, Test = 0.57



Random Forest Performance:
Accuracy: Train = 1.00, Test = 0.64
F1 Score: Train = 1.00, Test = 0.55
Precision: Train = 1.00, Test = 0.75
Recall: Train = 1.00, Test = 0.43



KNN Performance:
Accuracy: Train = 0.76, Test = 0.71
F1 Score: Train = 0.77, Test = 0.67
Precision: Train = 0.75, Test = 0.80
Recall: Train = 0.79, Test = 0.57




The optimal k value for KNN is: 9

KNN (Optimal) Performance:

Accuracy: Train = 0.71, Test = 0.86

F1 Score: Train = 0.72, Test = 0.83

Precision: Train = 0.68, Test = 1.00
Recall: Train = 0.76, Test = 0.71



c) Generate ROC plots for the training and testing datasets

```
# Import necessary libraries
import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_curve, auc

# Assuming the models (decision_tree, random_forest, knn) and datasets
(X_train_balanced, y_train_balanced, X_test, y_test) are already
defined

plt.figure(figsize=(10, 5))

# Plot PR curve for training data
plt.subplot(1, 2, 1)
plt.title('Precision-Recall Curve (Balanced Training Data)')
precision_dt, recall_dt, _ = precision_recall_curve(y_train_balanced,
decision_tree.predict_proba(X_train_balanced)[:, 1])
precision_rf, recall_rf, _ = precision_recall_curve(y_train_balanced,
random_forest.predict_proba(X_train_balanced)[:, 1])
precision_knn, recall_knn, _ =
```

```

precision_recall_curve(y_train_balanced,
knn.predict_proba(X_train_balanced)[: , 1])

plt.plot(recall_dt, precision_dt, label='Decision Tree (AUC =
 {:.2f})'.format(auc(recall_dt, precision_dt)))
plt.plot(recall_rf, precision_rf, label='Random Forest (AUC =
 {:.2f})'.format(auc(recall_rf, precision_rf)))
plt.plot(recall_knn, precision_knn, label='KNN (AUC =
 {:.2f})'.format(auc(recall_knn, precision_knn)))
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend()

# Plot PR curve for testing data
plt.subplot(1, 2, 2)
plt.title('Precision-Recall Curve (Testing Data)')
precision_dt_test, recall_dt_test, _ = precision_recall_curve(y_test,
decision_tree.predict_proba(X_test)[: , 1])
precision_rf_test, recall_rf_test, _ = precision_recall_curve(y_test,
random_forest.predict_proba(X_test)[: , 1])
precision_knn_test, recall_knn_test, _ =
precision_recall_curve(y_test, knn_optimal.predict_proba(X_test)[: ,
1])

plt.plot(recall_dt_test, precision_dt_test, label='Decision Tree (AUC
 = {:.2f})'.format(auc(recall_dt_test, precision_dt_test)))
plt.plot(recall_rf_test, precision_rf_test, label='Random Forest (AUC
 = {:.2f})'.format(auc(recall_rf_test, precision_rf_test)))
plt.plot(recall_knn_test, precision_knn_test, label='KNN (AUC =
 {:.2f})'.format(auc(recall_knn_test, precision_knn_test)))
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend()

plt.tight_layout()
plt.show()

```



Interpretation:

Training Data:

Decision Tree: AUC = 1.00 (perfect classification, potential overfitting)

Random Forest: AUC = 1.00 (perfect classification, potential overfitting)

KNN: AUC = 0.61 (good performance, less prone to overfitting)

Testing Data:

Decision Tree: AUC = 0.79 (significant drop in performance, suggests overfitting)

Random Forest: AUC = 0.76 (similar drop, indicates overfitting)

KNN: AUC = 0.94 (best performance, good generalization to unseen data)

The KNN model outperforms the others on the test data, indicating it generalizes better than the tree-based models, which exhibit signs of overfitting. Model overfitting occurs when the model captures noise in the training data, leading to poor performance on unseen data. This is likely due to the small dataset size, which emphasizes the need to increase the sample size for better representation. To mitigate overfitting, strategies such as increasing the dataset, applying cross-validation, adjusting model complexity, and using feature selection can be beneficial.